



Frist Pass Linux Kernel Dump

Dyi-Wu Liu

July, 2014

© Copyright 2013 Hewlett-Packard Development Company, L.P. The information contained herein is subject to change without notice.

Learning Objectives – Why we are here ?

Care Why

- Effective action plan development for outage case

Know Why

- Direct evidence finding instead of the best guess on action plan development

Know How

- Build it and test it

Know What


- What is kdump ?
- How does it work ?

QUINN, James Brian. "The intelligent enterprise a new paradigm." *The Executive*, 1992, 6.4: 48-63.



Agenda


- How kdump work
- Kdump setup/validation
- Serial console
- First Pass Kdump Analysis
- Labs



How kdump work

Kernel Crash Dumper

- Overview
- Kernel to Kernel Boot Loader
- User space tool and kernel system call
- Kexec on panic
- Work flow
- Pre-loading
- Post crash
- In core map
- ELF core file



How kdump work

Overview

- A kexec Based Kernel Crash Dumping Mechanism
- Dump capture from crashing kernel's context
- A new kernel, often called capture kernel, is booted after the crash
- Previous kernel's memory is preserved
- Dump is captured from the context of capture kernel
- Kernel to kernel boot loader enables booting a new kernel after a crash
- kexec is underlying kernel to kernel boot loader



How kdump work

Kernel to Kernel Boot Loader

- Running kernel acts as a loader for the new kernel
- System directly jumps from one kernel to another
 - Skips BIOS or Firmware stage
- Reboots are extremely fast
- Memory can be preserved across reboots
 - Since BIOS is skipped, it is left to the OS (capture kernel) to retain or erase memory



How kdump work

User space tool and kernel system call

- Allows a Linux kernel to boot another kernel
- Currently available on i386, x86_64, ia64 and ppc64 platforms
- Two components
 - User space tool – kexec(8)
 - Kernel System Call - kexec_load(2)

- Load a new kernel

```
kexec -l <kernel-image> --append=<options>
```

- Exec new kernel

```
kexec -e
```



How kdump work

Kexec on panic

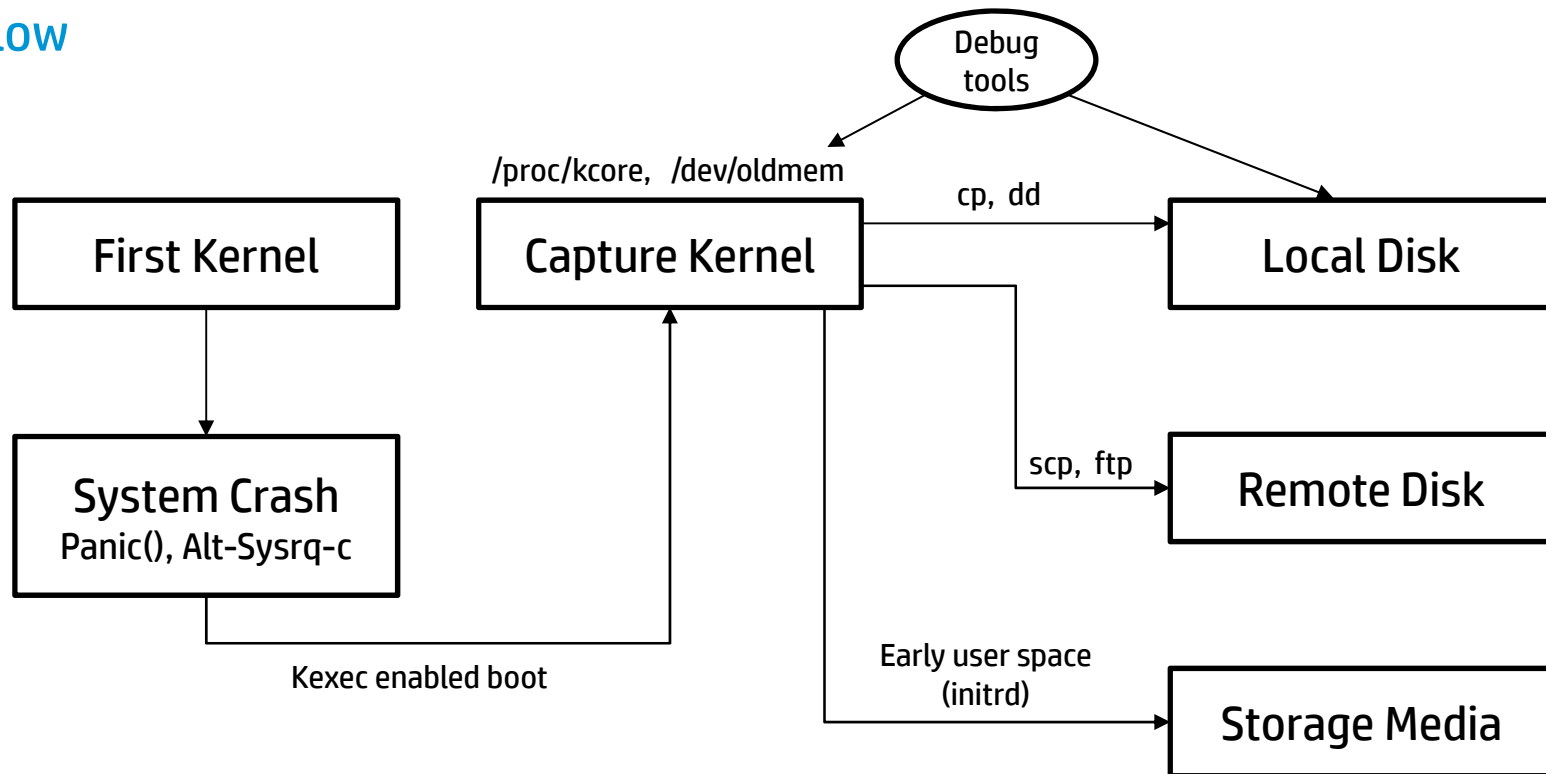
- Enables booting a new kernel after system crash
- Devices are not shutdown
- New kernel runs from a reserved memory location
 - Protection against ongoing DMA at the time of crash
- Loading capture kernel

```
kexec -p <kernel-image> --append=<options>
```
- Execution of capture kernel
 - panic()
 - Alt-Sysrq-c



How kdump work

Work flow



How kdump work


Pre-Loading

- Reserve memory for capture kernel (crashkernel=X@Y)
- Pre-load the capture kernel
- Capture kernel runs from reserved memory location
- For example (RHEL65)

```
/boot/grub/menu.lst
    crashkernel=auto

/proc/cmdline
    crashkernel=129M@0M

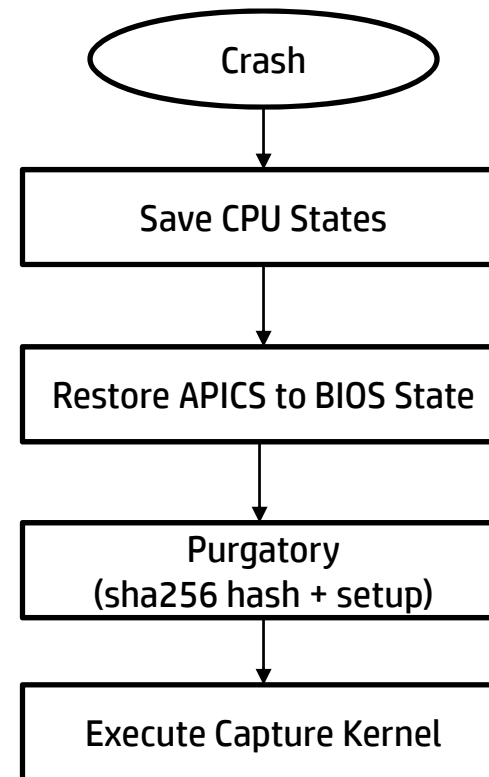
# grep Crash /proc/iomem
03000000-0b0fffff : Crash kernel
Y: 0x3000000 (48M)
X: 0xb0fffff-0x3000000+1 (129M)
```



How kdump work

Post Crash

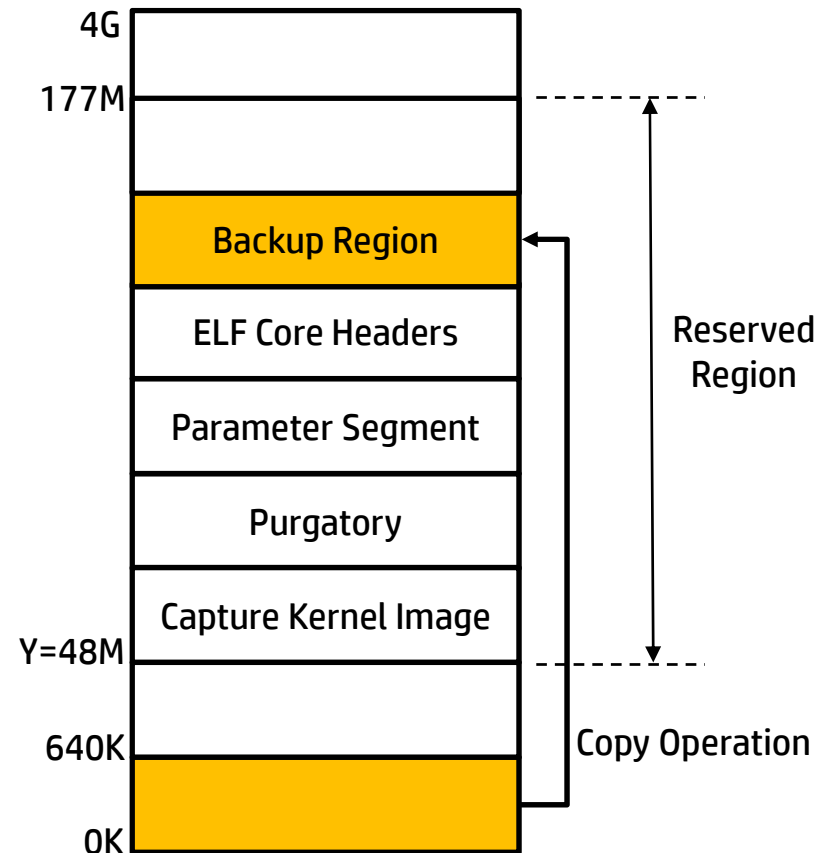
- CUP states are saved and other CPUs are halted
- CPU register states are saved on ELF note format
- 1K memory is reserved statically per CPU
- LAPIC/IOAPIC are disabled and put back into PIC or virtual wire mode
- Purgatory is run and control is transferred to capture kernel



How kdump work

In Core Map

- Dump information exchanged in an ELF format core file across the kernels (/proc/kcore)
- Kexec-tools prepare ELF headers and pre-load them in the reserved region
- Kexec-tools use /proc/iomem to retrieve system RAM information
- Content of first 640K of memory are backed up in backup region for the requirement to boot SMP capture kernel



How kdump work

ELF Core file

- Accessing dump image in ELF Core format
 - /proc/kcore
- Accessing dump image in linear raw format
 - /dev/oldmem
- ELF32/ELF64 format headers
- Physical addresses are filled for all the regions
- Virtual addresses are filled only for linearly mapped memory region

ELF Header	Program Header PT_NOTE	Program Header PT_NOTE	-----	Per CPU Register States	Dump Images
------------	------------------------	------------------------	-------	-------------------------	-------------




kdump setup/validation

Setup

- Required packages
- Magic SysRq key enable
- Configurations

Validation

- Check configuration after kdump enabled
- Force crash by Magic SysRq key c
- Force crash by Non-maskable interrupt (NMI)
- Saved vmcore



kdump setup/validation (RHEL 6.5)

Setup

- Required packages

```
# yum list installed | grep -e kexec -e kdump
```

```
kexec-tools.x86_64          2.0.0-273.el6      @anaconda-RedHatEnterpriseLinux-201311111358.x86_64/6.5
```

```
system-config-kdump.noarch 2.0.5-15.el6      @anaconda-RedHatEnterpriseLinux-201311111358.x86_64/6.5
```

```
# rpm -qa | grep -e kexec -e kdump
```

```
system-config-kdump-2.0.5-15.el6.noarch
```

```
kexec-tools-2.0.0-273.el6.x86_64
```

- Magic SysRq key enable

```
# grep sysrq /etc/sysctl.conf
```

```
kernel.sysrq = 1
```



kdump setup/validation (RHEL 6.5)

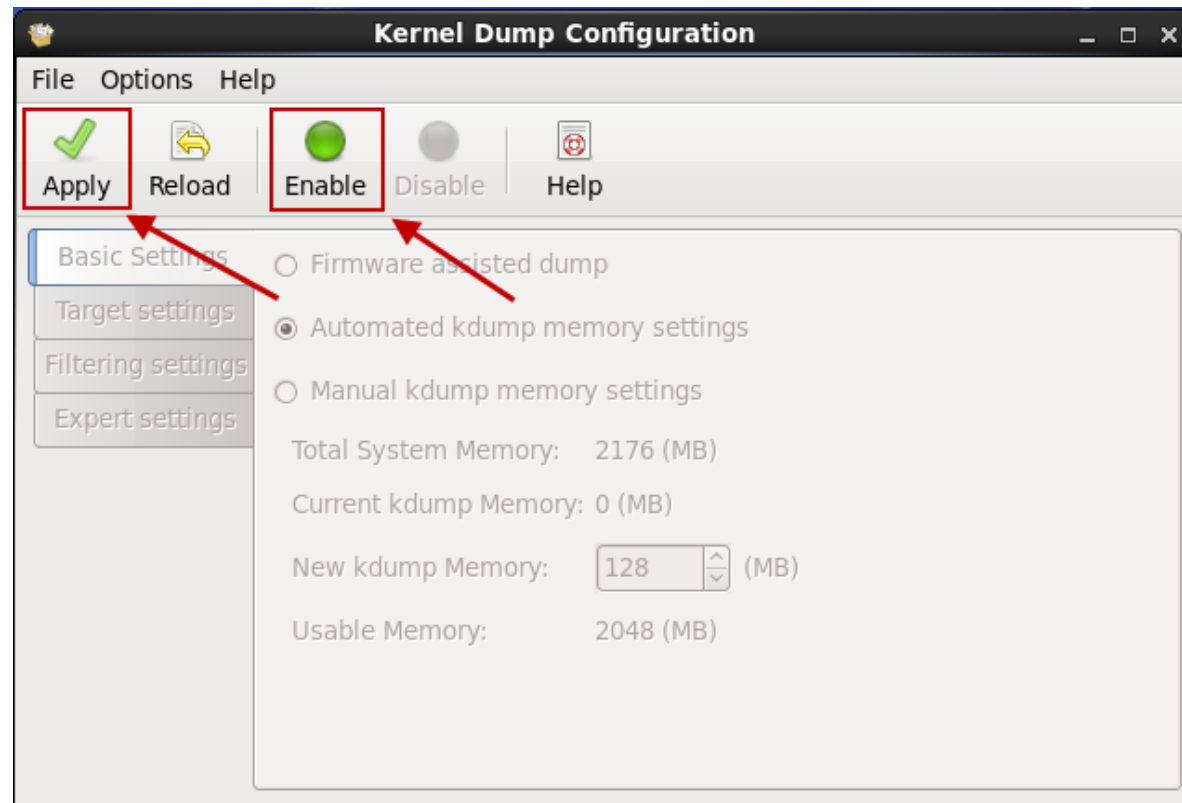
Setup

- Configurations

`system-config-kdump`

System reboot is required

- `kdump.conf(5)`
- `makedumpfile(8)`



kdump setup/validation (RHEL 6.5)

Validation

- Check configuration after kdump enabled

```
# grep crashkernel /boot/grub/menu.lst
        kernel /vmlinuz-2.6.32-431.el6.x86_64
root=/dev/mapper/VolGroup-lv_root ro rd_NO_LUKS
LANG=en_US.UTF-8 rd_NO_MD
rd_LVM_LV=VolGroup/lv_swap SYSFONT=latarcyrheb-
sun16 rd_LVM_LV=VolGroup/lv_root
KEYBOARDTYPE=pc KEYTABLE=us rd_NO_DM rhgb quiet
crashkernel=auto
```

```
# cat /proc/cmdline
root=/dev/mapper/VolGroup-lv_root ro rd_NO_LUKS
LANG=en_US.UTF-8 rd_NO_MD
rd_LVM_LV=VolGroup/lv_swap SYSFONT=latarcyrheb-
sun16 rd_LVM_LV=VolGroup/lv_root
KEYBOARDTYPE=pc KEYTABLE=us rd_NO_DM rhgb quiet
crashkernel=129M@0M
```

```
# chkconfig --list kdump
kdump 0:off 1:off 2:on 3:on 4:on 5:on 6:off
```

```
# service kdump status
Kdump is operational
```

```
# grep -v \# /etc/kdump.conf
path /var/crash
core_collector makedumpfile -c --message-level 1 -d 31
```

```
# sysctl kernel.sysrq
kernel.sysrq = 1
```



kdump setup/validation (RHEL 6.5)

Validation

- The c key of Magic SysRq perform a system crash

Keystroke : Alt + fn + SysRq + c

- Force crash by Magic SysRq key c

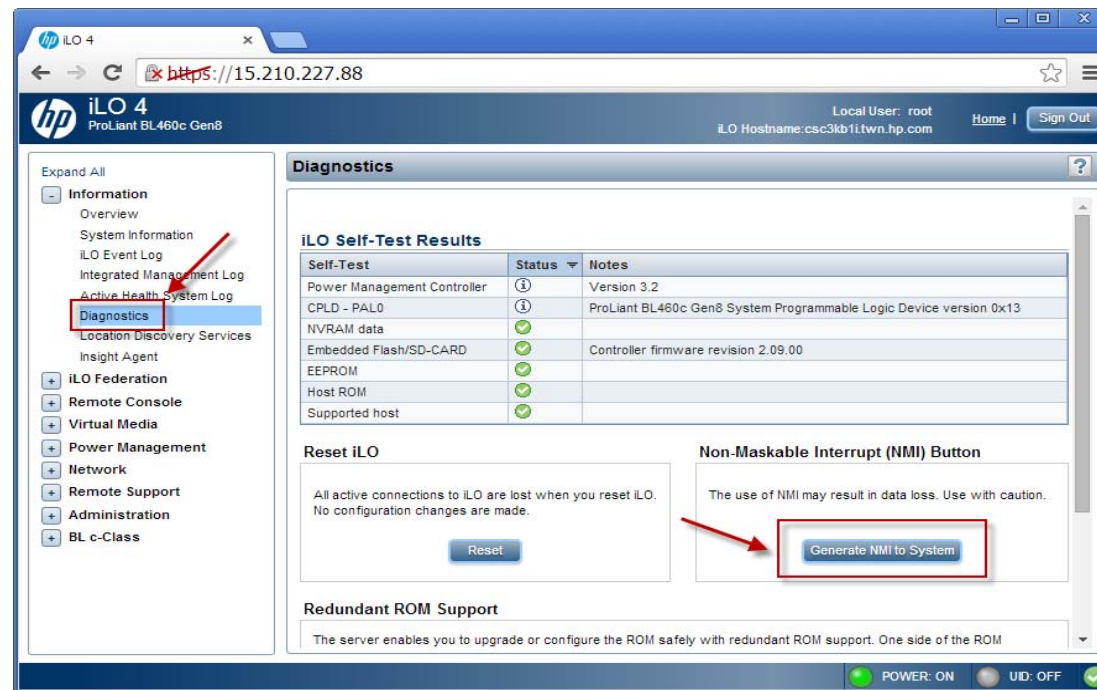
```
# cat /usr/local/bin/toc  
#!/bin/bash  
# echo 1 > /proc/sys/kernel/sysrq  
echo c > /proc/sysrq-trigger
```



kdump setup/validation (RHEL 6.5)

Validation

- Force crash by Non-maskable interrupt (NMI)



kdump setup/validation (RHEL 6.5)

Validation

- Saved vmcore

```
# ls -lR /var/crash/
/var/crash/:
total 4
drwxr-xr-x 2 root root 4096 Jul  7 12:40 127.0.0.1-2014-07-07-12:40:40

/var/crash/127.0.0.1-2014-07-07-12:40:40:
total 21448
-rw----- 1 root root 21868390 Jul  7 12:40 vmcore
-rw-r--r-- 1 root root    88328 Jul  7 12:40 vmcore-dmesg.txt

# cd /var/crash/127.0.0.1-2014-07-07-12\:40\:40/
# file *
vmcore:                                data
vmcore-dmesg.txt: ASCII English text, with very long lines

# strings vmcore | grep OSRELEASE
OSRELEASE=2.6.32-431.el6.x86_64

# uname -r
2.6.32-431.el6.x86_64
```



kdump setup/validation (SLES 11SP3)

Setup

- Required packages

```
# zypper se --installed-only | grep -e kdump -e kexec -e Summary -e -----
```

S	Name	Summary	Type
i	crash	Crash utility for live systems; netdump, diskdump, LKCD or mcore dumpfiles	package
i	kdump	Script for kdump	package
i	kexec-tools	Tools for fast kernel loading	package
i	yast2-kdump	Configuration of kdump	package

```
# rpm -qa | grep -e kdump -e kexec
```

```
kexec-tools-2.0.3-0.15.18  
yast2-kdump-2.17.26-0.8.17  
kdump-0.8.4-0.29.5
```

- Magic SysRq key enable

```
# tail -1 /etc/sysctl.conf
```

```
kernel.sysrq = 1
```

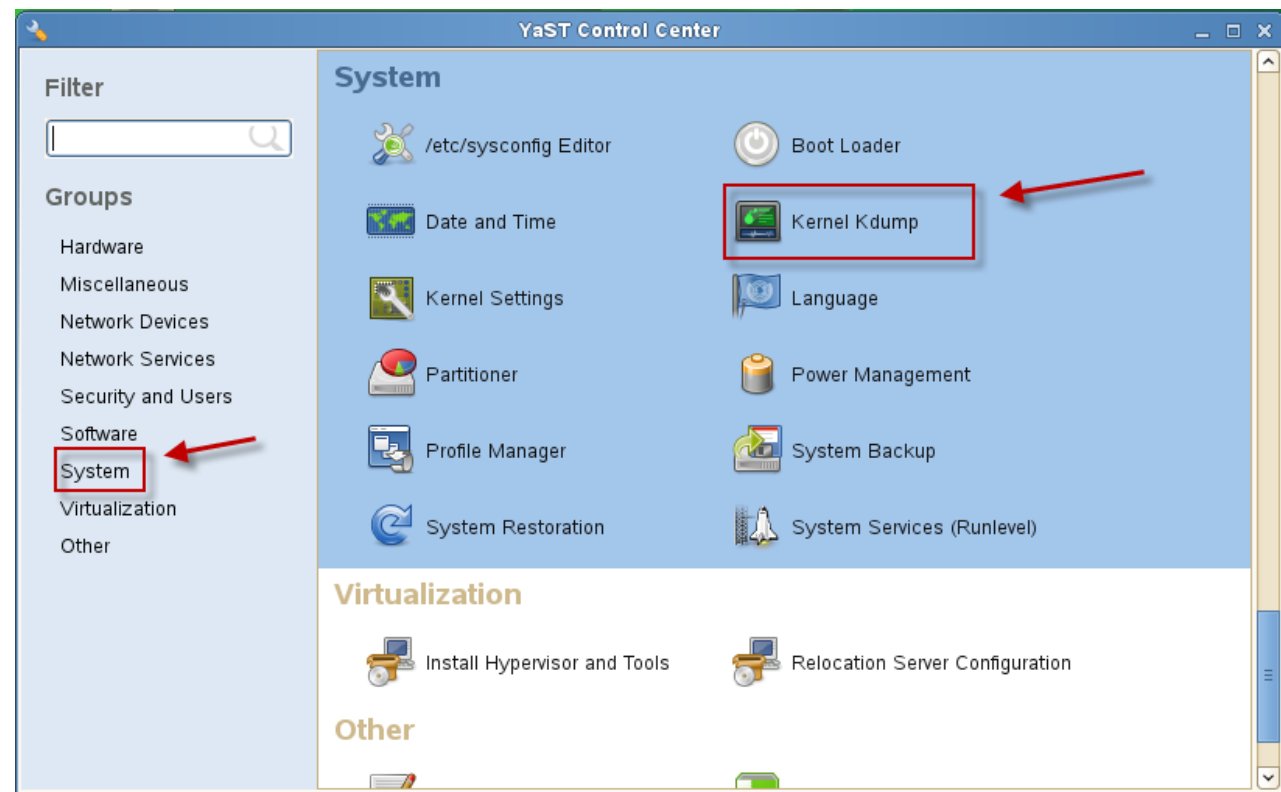


kdump setup/validation (SLES 11SP3)

Setup

- Configurations

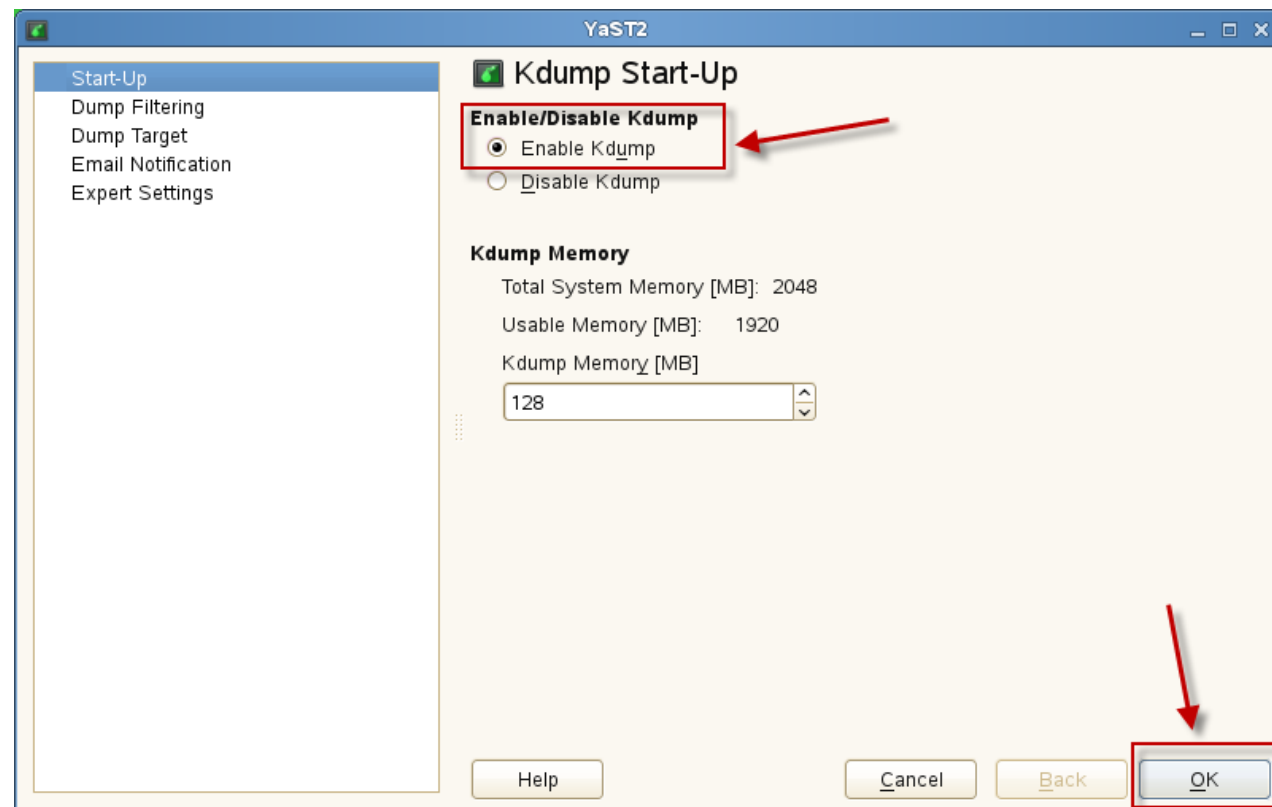
yast2



kdump setup/validation (SLES 11SP3)

Setup

- Configurations
- **System reboot is required**
- kdump.conf(5)
- makedumpfile(8)



kdump setup/validation (SLES 11SP3)

Validation

- Check configuration after kdump enabled

```
# grep crashkernel /boot/grub/menu.lst
        kernel /boot/vmlinuz-3.0.76-0.11-default
root=/dev/sda2 resume=/dev/sda1 splash=silent
showopts crashkernel=256M-:128M vga=0x314
```

```
# cat /proc/cmdline
root=/dev/sda2 resume=/dev/sda1 splash=silent
crashkernel=256M-:128M vga=0x314
```

```
# chkconfig --list boot.kdump
boot.kdump          0:off 1:off 2:off
3:off 4:off 5:off 6:off B:on
```

```
# service boot.kdump status
kdump kernel loaded          running
```

```
# grep -v \# /etc/sysconfig/kdump | grep -v ^$ | grep -v \"\
KDUMP_CPUS="1"
KDUMP_IMMEDIATE_REBOOT="yes"
KDUMP_SAVEDIR="file:///var/crash"
KDUMP_KEEP_OLD_DUMPS="5"
KDUMP_FREE_DISK_SIZE="64"
KDUMP_VERBOSE="3"
KDUMP_DUMPLEVEL="31"
KDUMP_DUMPFORMAT="lzo"
KDUMP_CONTINUE_ON_ERROR="true"
KDUMP_COPY_KERNEL="yes"
KDUMP_NETCONFIG="auto"

# sysctl kernel.sysrq
kernel.sysrq = 1
```



kdump setup/validation (SLES 11SP3)

Validation

- The c key of Magic SysRq perform a system crash

Keystroke : Alt + fn + SysRq + c

- Force crash by Magic SysRq key c

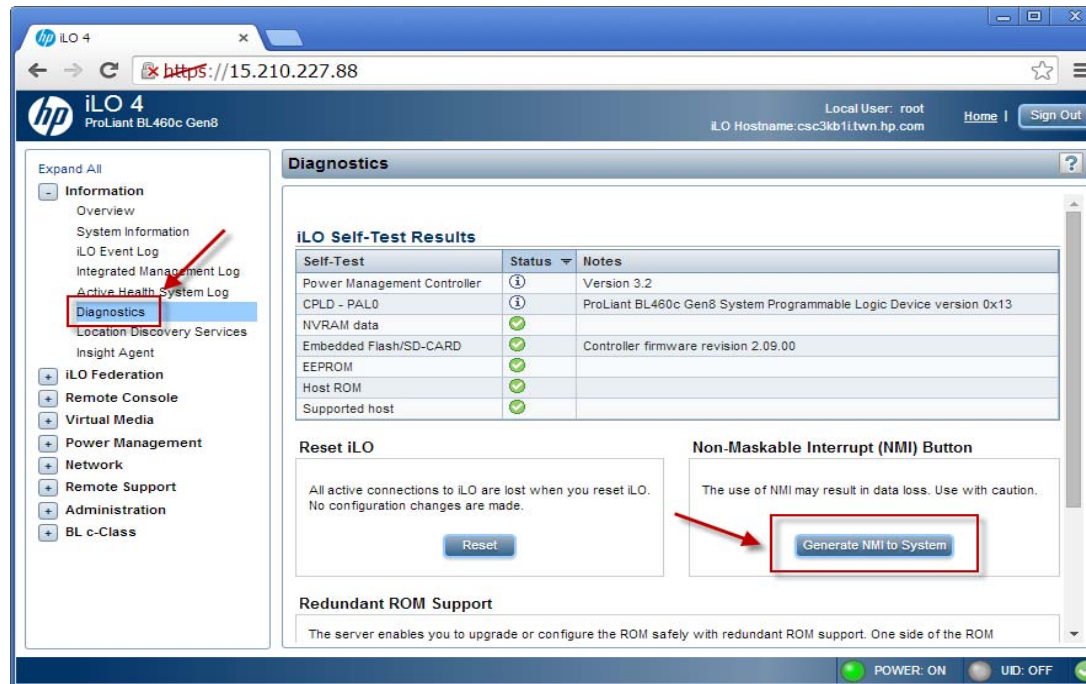
```
# cat /usr/local/bin/toc
#!/bin/bash
# echo 1 > /proc/sys/kernel/sysrq
echo c > /proc/sysrq-trigger
```



kdump setup/validation (SLES 11SP3)

Validation

- Force crash by Non-maskable interrupt (NMI)



kdump setup/validation (SLES 11SP3)

Validation

- Saved vmcore

```
# ls -lR /var/crash/
/var/crash/:
total 4
drwxr-xr-x 2 root root 4096 Jul  6 08:14 2014-07-06-00:32
/var/crash/2014-07-06-00:32:
total 36460
-rw----- 1 root root  112018 Jul  6 00:32 dmesg.txt
-rw-r--r-- 1 root root    179 Jul  6 00:32 README.txt
-rw-r--r-- 1 root root 2067870 Jul  6 00:32 System.map-3.0.76-0.11-default
-rw----- 1 root root 30481332 Jul  6 00:32 vmcore
-rw-r--r-- 1 root root 4605902 Jul  6 00:32 vmlinux-3.0.76-0.11-default.gz

# cd /var/crash/2014-07-06-00\:32/
# file *
dmesg.txt:          ASCII C++ program text
README.txt:         ASCII text
System.map-3.0.76-0.11-default: ASCII text
vmcore:            data
vmlinux-3.0.76-0.11-default.gz: gzip compressed data, from Unix, max compression

# strings vmcore | grep OSRELEASE|head -1
OSRELEASE=3.0.76-0.11-default

# uname -r
3.0.76-0.11-default
```



Serial console

Serial Console on Bare Metal

- BIOS configuration
- GRUB configuration
- HP iLO VSP (Virtual Serial Port)
- Connect through iLO

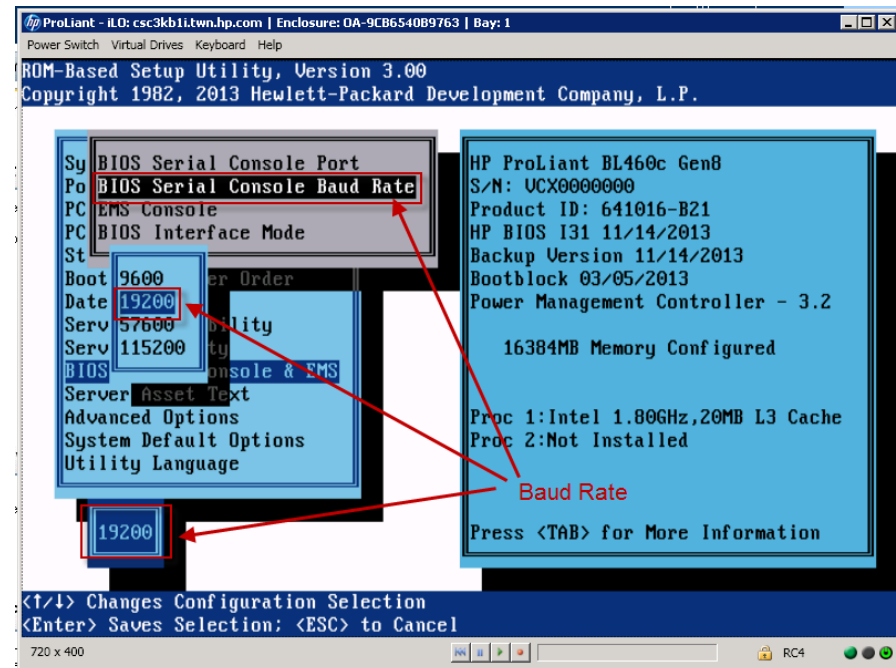
Serial Console on VM Guest

- BIOS configuration
- GRUB configuration
- Serial port on VM Guest
- Connect through named pipe via putty



- BIOS Serial Console to COM2
- BIOS Serial Console Baud Rate to 19200

-
- HP ProLiant - iLO: csc3kbl1.twn.hp.com | Enclosure: OA-9C86540B9763 | Bay: 1
- Power Switch Virtual Drives Keyboard Help
- ROM-Based Setup Utility, Version 3.00
- Copyright 1982, 2013 Hewlett-Packard Development Company, L.P.
- System BIOS Serial Console Port
- Port BIOS Serial Console Baud Rate
- PC EMS Console
- PC BIOS
- Setup Auto
- Boot Disabled
- Date COM 1; IRQ4; IO: 3F8h-3FFh
- Service COM 2; IRQ3; IO: 2F8h-2FFh
- Service
- BIOS Serial Console & EMS
- Server Asset Text
- Advanced Options
- System Default Options
- Utility Language
- COM 2; IRQ3; IO: 2F8h-2FFh
- HP ProLiant BL460c Gen8
- S/N: UCX0000000
- Product ID: 641016-B21
- HP BIOS I31 11/14/2013
- Backup Version 11/14/2013
- Bootblock 03/05/2013
- Power Management Controller - 3.2
- 16384MB Memory Configured
- Proc 1: Intel 1.80GHz, 20MB L3 Cache
- Proc 2: Not Installed
- BIOS Serial Console Port
- Press <TAB> for More Information
- <F1> Changes Configuration Selection
- <Enter> Saves Selection; <ESC> to Cancel
- 720 x 400
- RC4



Serial Console - Bare Metal (RHEL 6.5)

Name of COM2, IRQ and I/O Port

```
# setserial -g /dev/ttyS[01]
/dev/ttyS0, UART: 16550A, Port: 0x03f8, IRQ: 4
/dev/ttyS1, UART: 16550A, Port: 0x02f8, IRQ: 3
```

Kernel line in /boot/grub/menu.lst

```
#boot=/dev/sda
default=0
timeout=5
splashimage=(hd0,0)/grub/splash.xpm.gz
serial --unit=1 --speed=19200 --word=8 --parity=no --stop=1
terminal --timeout=5 serial console
hiddenmenu
title Red Hat Enterprise Linux (2.6.32-431.el6.x86_64)
    root (hd0,0)
    kernel /vmlinuz-2.6.32-431.el6.x86_64 ro root=/dev/mapper/vg00-root rd_NO_LUKS
rd_LVM_LV=vg00/swap LANG=en_US.UTF-8 rd_LVM_LV=vg00/root rd_NO_MD SYSFONT=latacyrheb-sun16
crashkernel=128M KEYBOARDTYPE=pc KEYTABLE=us rd_NO_DM rhgb console=tty0 console=ttyS1,19200n8
    initrd /initramfs-2.6.32-431.el6.x86_64.img
```



Serial Console - Bare Metal (RHEL 6.5)

Connect through iLO

```
# ssh csc3kbli
root@csc3kbli's password:
User:root logged-in to csc3kbli.twn.hp.com(15.210.227.88 / fe80::a65d:36ff:fefe:bb9)
iLO 4 Advanced for BladeSystem 1.32 at Nov 05 2013
Server Name:
Server Power: On
</>hpiLO-> vsp

Virtual Serial Port Active: COM2
Starting virtual serial port.
Press 'ESC (' to return to the CLI Session.

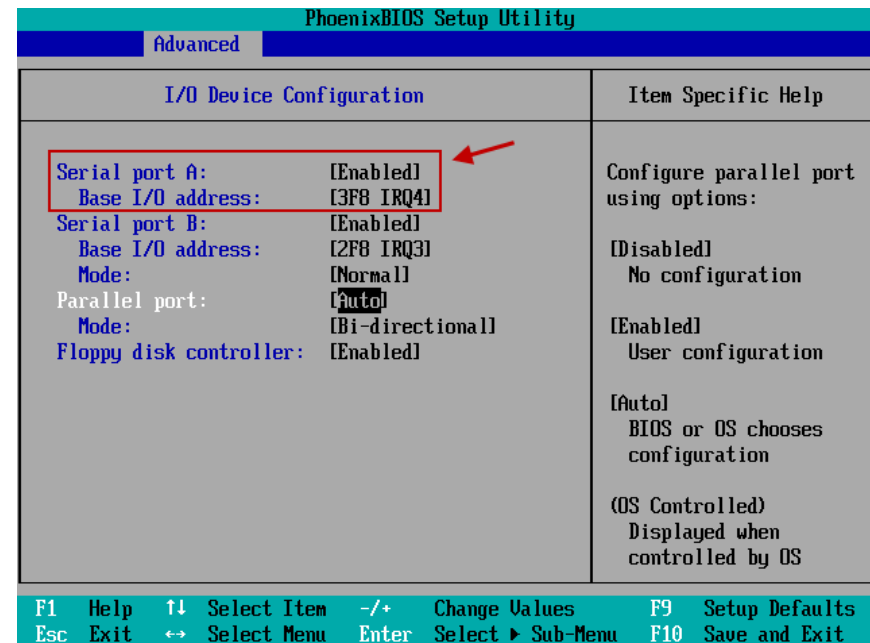
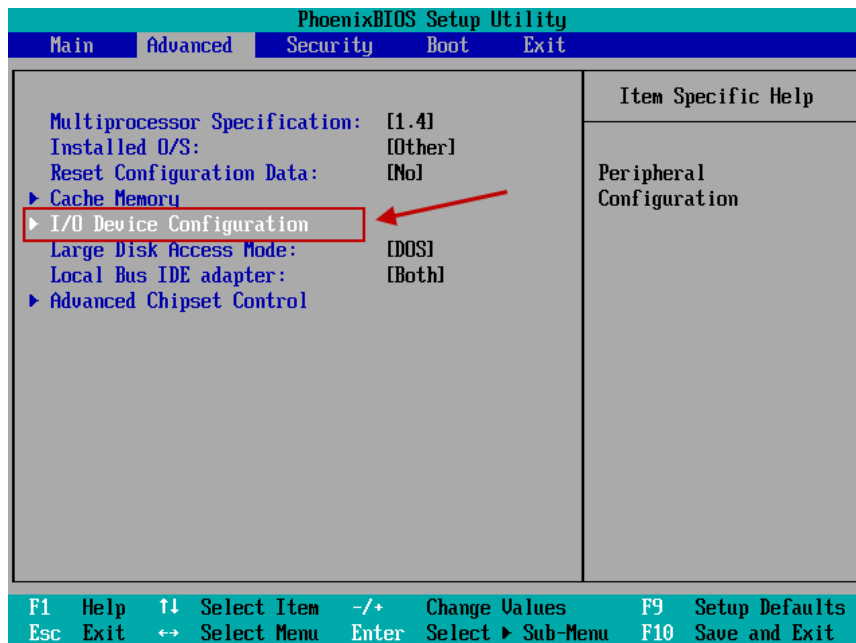
Red Hat Enterprise Linux Server release 6.5 (Santiago)
Kernel 2.6.32-431.el6.x86_64 on an x86_64
csc3kb1.twn.hp.com login:
```

Console session log in text



Serial Console – VM Guest (RHEL 6.5)

- Enable serial port A
- Base I/O address, IRQ number



Serial Console – VM Guest (RHEL 6.5)

Name of port A, IRQ and I/O Port

```
# setserial -g /dev/ttyS[01]
/dev/ttyS0, UART: 16550A, Port: 0x03f8, IRQ: 4
/dev/ttyS1, UART: 16550A, Port: 0x02f8, IRQ: 3
```

Kernel line in /boot/grub/menu.lst

```
#boot=/dev/sda
default=0
timeout=5
splashimage=(hd0,0)/grub/splash.xpm.gz
serial --unit=0 --speed=19200 --word=8 --parity=no --stop=1
terminal --timeout=5 serial console
# hiddenmenu
title Red Hat Enterprise Linux (2.6.32-431.el6.x86_64)
    root (hd0,0)
    kernel /vmlinuz-2.6.32-431.el6.x86_64 ro root=/dev/mapper/vg00-root rd_NO_LUKS
rd_LVM_LV=vg00/swap LANG=en_US.UTF-8 rd_LVM_LV=vg00/root rd_NO_MD SYSFONT=latarcyrheb-sun16
crashkernel=128M KEYBOARDTYPE=pc KEYTABLE=us rd_NO_DM rhgb console=tty0 console=ttyS0,19200n8
    initrd /initramfs-2.6.32-431.el6.x86_64.img
```



Serial Console – VM Guest (RHEL 6.5)

Enable ttyS0 and console device in /etc/securetty

```
# grep -e console -e ttyS0 /etc/securetty
ttyS0
console
```

Check the agetty process

```
# initctl status serial DEV=ttyS0
serial (ttyS0) start/running, process 2321
# ps -ef | grep agetty | grep -v grep
root      2321      1  0 13:39 ttyS0      00:00:00 /sbin/agetty /dev/ttyS0 19200 vt100-navgrep console
```

Connect through putty

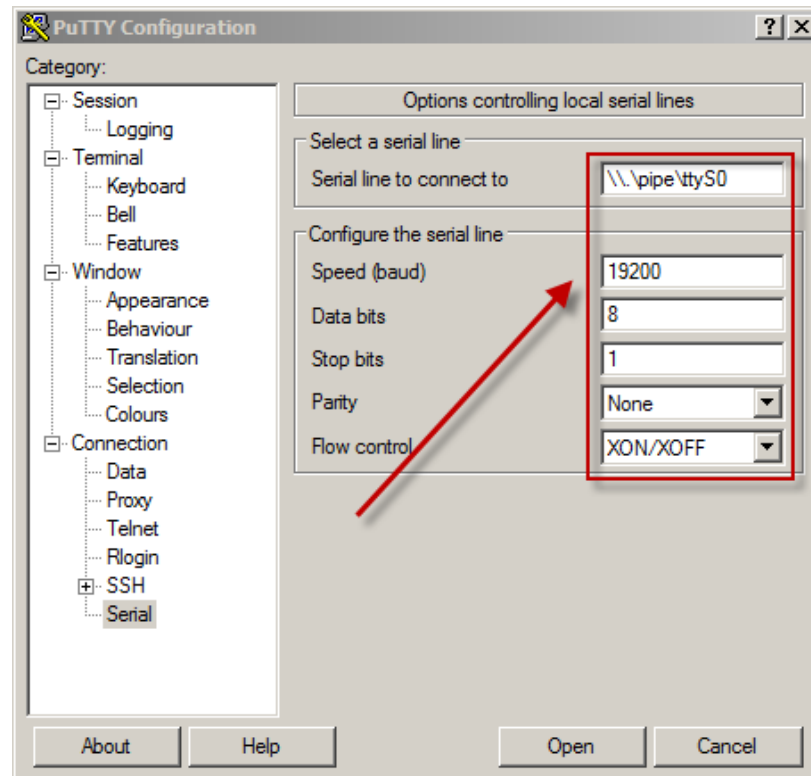
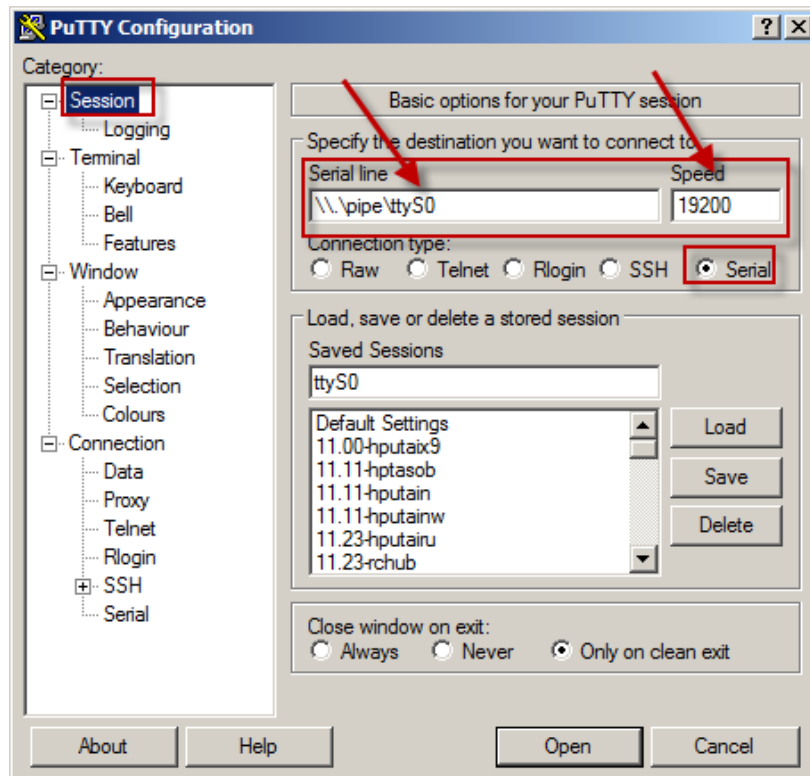
```
Red Hat Enterprise Linux Server release 6.5 (Santiago)
Kernel 2.6.32-431.el6.x86_64 on an x86_64
```

```
linux3.asiapacific.hpqcorp.net login: root
Password:
Last login: Mon Jul  7 13:39:37 from kcz.twn.hp.com
# tty
/dev/ttyS0
```



Serial Console – putty client (RHEL 6.5)

Named pipe connection from putty



Serial Console – VM Guest (SLES 11 SP3)

Name of port A, IRQ and I/O Port

```
# setserial -g /dev/ttyS[01]
/dev/ttyS0, UART: 16550A, Port: 0x03f8, IRQ: 4
/dev/ttyS1, UART: 16550A, Port: 0x02f8, IRQ: 3
```

Port A entry in /etc/init.d/setserial

```
# grep 0x3F8 /etc/init.d/setserial
run_setserial /dev/ttyS0 uart 16550A port 0x3F8 irq 4
```

Kernel line in /boot/grub/menu.lst

```
##YaST - generic_mbr
# gfxmenu (hd0,1)/boot/message
serial --unit=0 --speed=19200
terminal --timeout=5 serial console
##YaST - activate

###Don't change this comment - YaST2 identifier: Original name: linux###
title SUSE Linux Enterprise Server 11 SP3 - 3.0.76-0.11
    root (hd0,1)
    kernel /boot/vmlinuz-3.0.76-0.11-default root=/dev/sda2 resume=/dev/sda1 splash=silent showopts
    crashkernel=256M-:128M vga=0x314 console=tty0 console=ttyS0,19200
    initrd /boot/initrd-3.0.76-0.11-default
```



Serial Console – VM Guest (SLES 11 SP3)

Enable ttyS0 and console device in /etc/securetty

```
# tail -2 /etc/securetty
ttyS0
console
```

Enableagetty in /etc/inittab

```
# grep console /etc/inittab
cons:12345:respawn:/sbin/smartagetty -L 19200 console
```

Connect trough putty



```
Welcome to SUSE Linux Enterprise Server 11 SP3 (x86_64) - Kernel 3.0.76-0.11-default (console).
```

```
linux-c4o5 login: root
Password:
Last login: Sun Jul  6 23:50:56 CST 2014 on console
# tty
/dev/console
```



Serial Console – VM Guest (SLES 11 SP3)

Named pipe connection from putty



First Pass Kdump Analysis

Types of hangs

Type of crashes

Vmcore analysis tools

- crash64
- mpykdump64.so
- crashinfo
- Remote crash
 - ✓ crash-server
 - ✓ crash-client – text version
 - ✓ crash-client – web version



Moment of ecstasy-

Doc Debugalov finds
the elusive bug corrupting
the heap...

First Pass Kdump Analysis

Type of hangs

- Network hang
 - loss of network connectivity but system still running
 - Does not respond to a ping
 - Users are unable to access system with ssh
 - The system can be accessed via the iLO, or serial port
 - Partial network hanging? check switches / routers
- I/O hang
 - Some storage has failed and the system is waiting for it
 - System responds to ping
 - Users hang after connecting
- Complete hang
 - System unresponsive to network and console



First Pass Kdump Analysis

Type of crashes

- **Oops**
 - One process gets killed immediate
 - A subsystem is left in an unknown state
- **Panic**
 - Something happens in the kernel that it is unable to continue from
- **BUG**
 - Macro in the kernel testing for a condition that should not happen
- **OOM Killer**
 - Kernel memory starvation where the kernel must kill processes to keep
- **NMI**
 - Hardware generated or virtual button on the iLO
- **Other**
 - The perhaps overlooked loss of power causing an unexplained system outage



First Pass Kdump Analysis

Vmcore analysis tools

- crash64
- mpykdump64.so
- crashinfo
- crash-server
- crash-client – text version
- crash-client – web version
- <ftp://kczb.twn.hp.com/dist/debuginfo/tools/>
- <http://sourceforge.net/projects/pykdump/>



First Pass Kdump Analysis

crash64

```
# crash64
crash64 7.0.3
crash64: cannot find booted kernel -- please enter namelist argument
Usage:
  crash [OPTION]... NAMELIST MEMORY-IMAGE  (dumpfile form)
  crash [OPTION]... [NAMELIST]                (live system form)
Enter "crash64 -h" for details.
```

- NAMELIST

- The debuginfo file from RedHat for RHEL or from Novell for SLES
- <http://bl460node3.alf.cpqcorp.net/debuginfo/>
- <ftp://kczb.twn.hp.com/dist/debuginfo>

- MEMORY-IMAGE

- Saved vmcore from kdump
- Live system /dev/crash



First Pass Kdump Analysis

crash64 invocation on RHEL

```
# ll
total 152036
-rw----- 1 root root 21868390 Jul  7 12:40 vmcore
-rw-r--r-- 1 root root 88328 Jul  7 12:40 vmcore-dmesg.txt
-rw-r--r-- 1 root root 133715563 Jun 19 14:33 vmlinux-2.6.32-431.el6.x86_64

# crash64 vmlinux-2.6.32-431.el6.x86_64 vmcore
crash64 7.0.3
Copyright (C) 2002-2013 Red Hat, Inc.
Snipped ...
crash64> help
*
alias      foreach   mach      repeat    timer
ascii      fuser     mod       runq      tree
bt         gdb       mount     search    union
btop       help      net       set       vm
dev        ipcs      p         sig       vtop
dis        irq       ps        struct    waitq
eval       kmem      ptob      swap      whatis
exit       list      ptov      sym       wr
extend     log       rd        task      q
```



First Pass Kdump Analysis

crash64 invocation on SLES

```
# crash64 -s vmlinux-3.0.76-0.11-default.debug vmcore
crash64: vmlinux-3.0.76-0.11-default.debug: no text and data contents

crash64: The namelist argument supplied in this case is a debuginfo file,
which must be accompanied by the kernel file from which it was derived.

# ll
total 147192
-rw----- 1 root root 112018 Jul 6 00:32 dmesg.txt
-rw-r--r-- 1 root root 179 Jul 6 00:32 README.txt
-rw-r--r-- 1 root root 2067870 Jul 6 00:32 System.map-3.0.76-0.11-default
-rw----- 1 root root 30481332 Jul 6 00:32 vmcore
-rw-r--r-- 1 root root 22701625 Jul 6 00:32 vmlinux-3.0.76-0.11-default
-rw-r--r-- 1 root root 95177924 Jul 7 00:30 vmlinux-3.0.76-0.11-default.debug

# crash64 -s vmlinux-3.0.76-0.11-default vmcore
crash64> sys
      KERNEL: vmlinux-3.0.76-0.11-default
      DUMPFILE: vmcore [PARTIAL DUMP]
      CPUS: 2
      DATE: Sun Jul 6 00:32:10 2014
Snipped ...
```



First Pass Kdump Analysis

crashinfo – mpykdump64.so

```
# ll /usr/local/lib
total 3068
-rwxr-xr-x 1 root root 3140837 Nov 16 2013 mpykdump64.so

# crash64 -s vmlinux-2.6.32-431.el6.x86_64 vmcore
crash64> extend /usr/local/lib/mpykdump64.so
/usr/local/lib/mpykdump64.so: shared object loaded
crash64> help
*          extend          mach          repeat          timer
alias      files          mod          runq           tree
ascii      foreach        mount        search          union
bt         fuser          net          set            vm
btop       gdb             nfsshow      sig            vtop
crashinfo  help             p           struct          waitq
dev        ipcs            ps          swap           whatis
dis        irq             pte         sym            wr
epython    knem             ptob        sys            xportshow
eval       list             ptov        task           q
exit       log             rd          taskinfo

crash64> crashinfo -v > crashinfo.txt
crash64> crashinfo | grep Dump
>-----| How This Dump Has Been Created |-----<
Dump has been initiated: with sysrq
```



First Pass Kdump Analysis

Command file -i myfirstpass

```
# cat myfirstpass
sys
sys -c
kmem -i
kmem -f
bt
bt -a
foreach bt | grep spin
ps | grep -v " IN "
runq | grep PID
mod
log
quit

# crash64 -s -i myfirstpass vmlinux-2.6.32-431.el6.x86_64 vmcore > my.txt
```

Have fun with my.txt !!!



First Pass Kdump Analysis

Remote crash

- Time consuming on two transfers on the same vmcore

```
Customer  ---> ftp.usa.hp.com ---> ERT
vmcore      vmcore      vmcore
```

- Install debuginfo rpm and crash-server on customer site
- Open one port, ie tcp/5111 port
- Run crash-client and look at the vmcore remotely
- No vmcore transfer

```
crash-server <--- tcp/5111 ---> crash-client
```

Remote crash components

- Crash-server
- Crash-client – text version
- Crash-client – web version



First Pass Kdump Analysis

Crash-server

- To be run on the system where vmlinux and vmcore resides
- Reads input (commands) from and writes output (result) to a socket
- Listens on one port between 5111 and 5130
- 20 instances can be run on a single system
- Versions
 - crash-server32 (x86)
 - crash-server64 (x86_64)
 - crash-serveria64 (ia64)
- Produces a log file /tmp/wtec-log-username-port#



First Pass Kdump Analysis

Crash-server

```
# pwd
/var/crash/127.0.0.1-2014-07-07-12:40:40
# ll
total 152040
-rw-r--r-- 1 root root      108 Jul  7 16:30 myfirstpass
-rw----- 1 root root 21868390 Jul  7 12:40 vmcore
-rw-r--r-- 1 root root   88328 Jul  7 12:40 vmcore-dmesg.txt
-rw-r--r-- 1 root root 133715563 Jun 19 14:33 vmlinux-2.6.32-431.el6.x86_64

# crash-server64 -s vmlinux-2.6.32-431.el6.x86_64 vmcore
7ff6f000crash-server ready. Listening on port # 5111
All commands will be recorder on file /tmp/wtec-log-root-5111
The web client can be started by going to:
http://linuxdb.corp.hp.com/crash-client/main.php?remote\_host=linuxbox.abc.com&port=5111
Processing sys ....Done
Processing bt ....Done
Processing extend /usr/local/lib/mpykdump64.so ....Done
Processing crashinfo ....Done
```



First Pass Kdump Analysis

Crash-client – text version

- Communicates with server via a socket (tcp/5111 ~ tcp/5130)
- Basic and limited tty (text) app
- Loops prompting for crash commands
- Works for all regular crash commands
- Supports pipes | and redirections >
- Support readline and history



First Pass Kdump Analysis

Crash-client – text version

```
# crash-client64 linuxbox.abc.com 5111
Enter file name to log crash output: /tmp/abc.log
crash-client Version 1.0
To report bugs or enhancements please email
a detail report stating the problem and vmcore location to:
eddie.quinteros@hp.com

      KERNEL: vmlinux-2.6.32-431.el6.x86_64
      DUMPFILE: vmcore [PARTIAL DUMP]
      CPUS: 2
      DATE: Mon Jul  7 12:40:37 2014
      UPTIME: 00:15:36
LOAD AVERAGE: 0.00, 0.00, 0.00
      TASKS: 212
      NODENAME: linuxbox.abc.com
      RELEASE: 2.6.32-431.el6.x86_64
      VERSION: #1 SMP Sun Nov 10 22:19:54 EST 2013
      MACHINE: x86_64 (2494 Mhz)
      MEMORY: 2 GB
      PANIC: "Oops: 0002 [#1] SMP " (check log for details)

crash-client> sys
crash-client> bt
crash-client> extend /usr/local/lib/mpykdump64.so
crash-client> crashinfo > crashinfo.txt
```



First Pass Kdump Analysis

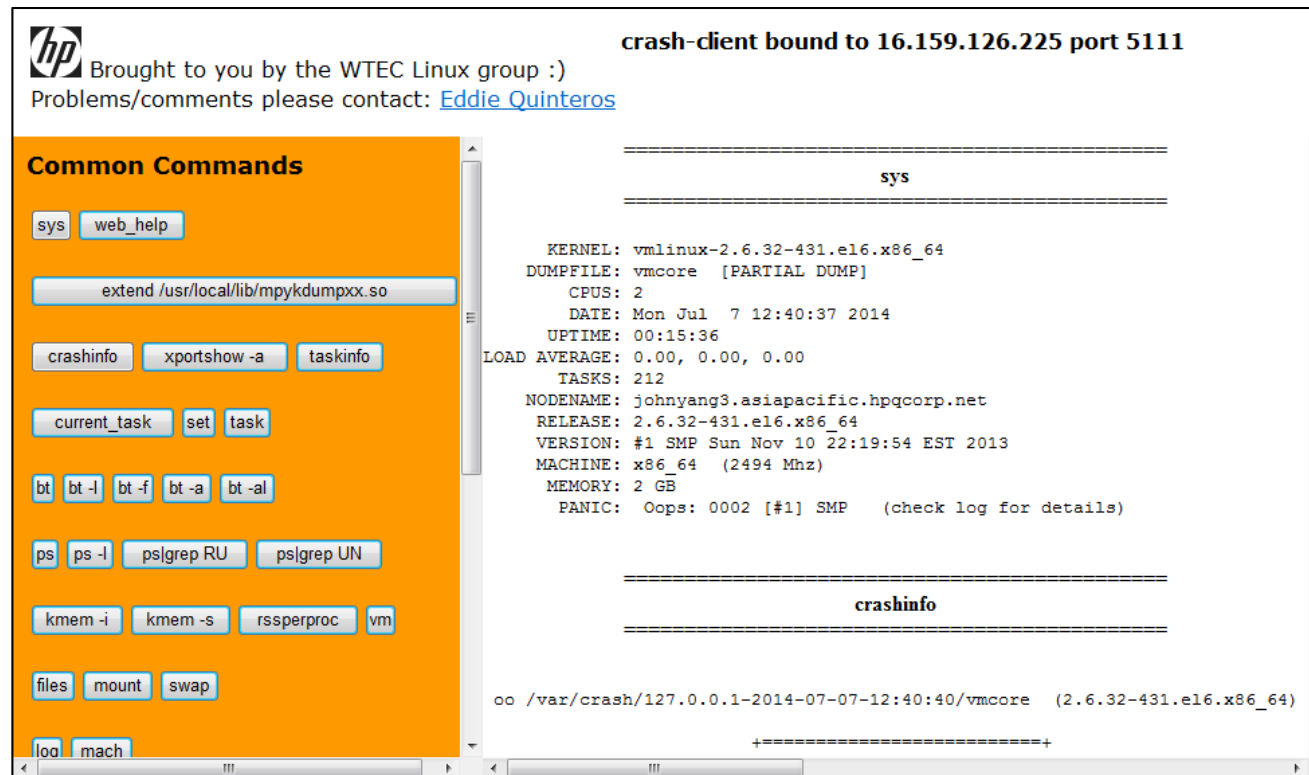
Crash-client – web version

- More advanced than the text version
- Easier to read vmcore by producing links on most command
- Tied to linuxdb.corp.hp.com
- Links to functions and source code lines on linuxdb.corp.hp.com
- It can be expanded to include customer commands and links



First Pass Kdump Analysis

- http://linuxdb.corp.hp.com/crash-client/main.php?remote_host=16.159.126.225&port=5111



The screenshot displays the 'crash-client' web interface. At the top left is the HP logo and the text 'Brought to you by the WTEC Linux group :)'. Below this is a link to 'Eddie Quinteros' for problems/comments. The main title is 'crash-client bound to 16.159.126.225 port 5111'. The interface is divided into two main sections. On the left is an orange sidebar titled 'Common Commands' containing various buttons for system and crash-related actions. On the right is a white area displaying system information under the heading 'sys'. The system information includes kernel version, dumpfile, CPU count, date, uptime, load average, tasks, nodename, release, version, machine type, memory, and panic details. Below this is a section titled 'crashinfo' which shows the path to the crash dump file.

Common Commands

Buttons: sys, web_help, extend /usr/local/lib/mpykdumpxx.so, crashinfo, xportshow -a, taskinfo, current_task, set, task, bt, bt -l, bt -f, bt -a, bt -al, ps, ps -l, pslgrep RU, pslgrep UN, kmem -i, kmem -s, rssperproc, vm, files, mount, swap, loc, mach

crash-client bound to 16.159.126.225 port 5111

Brought to you by the WTEC Linux group :)
Problems/comments please contact: [Eddie Quinteros](#)

sys

KERNEL: vmlinux-2.6.32-431.el6.x86_64
DUMPFILE: vmcore [PARTIAL DUMP]
CPUS: 2
DATE: Mon Jul 7 12:40:37 2014
UPTIME: 00:15:36
LOAD AVERAGE: 0.00, 0.00, 0.00
TASKS: 212
NODENAME: johnyang3.asiapacific.hpqcorp.net
RELEASE: 2.6.32-431.el6.x86_64
VERSION: #1 SMP Sun Nov 10 22:19:54 EST 2013
MACHINE: x86_64 (2494 Mhz)
MEMORY: 2 GB
PANIC: Oops: 0002 [#1] SMP (check log for details)

crashinfo

oo /var/crash/127.0.0.1-2014-07-07-12:40:40/vmcore (2.6.32-431.el6.x86_64)



Labs

Exercises

- kdump setup and validation
- Enable the serial console
- vmcore analysis
 - crash64
 - crashinfo
 - crash-server
 - crash-client – text version
 - crash-client – web version



Labs (cont.)

RedHat 6.5 guest VM

- VMWare Workstation 9.0 virtual machine
- SELinux, Ipv4/ipv6 packet filter disabled
- Yum repository on ftp://127.0.0.1/dist/linux/RedHat/65Server/os/x86_64

SLES 11 SP3 guest VM

- VMWare Workstation 9.0 virtual machine
- SELinux, Ipv4/ipv6 packet filter disabled
- Zypper repository on <ftp://127.0.0.1/dist/linux/SLES/sles11sp3>

Tools

- ERT kernels source code Page - <http://linuxdb.corp.hp.com/>
- debuginfo repository on <ftp://kczb.twn.hp.com/dist/debuginfo>



AHA !!!



debug..."



Learning Objectives Review - Take away

Care Why

- Effective action plan development for outage case
- With the vmcore on hand, do you feel more confidence on cause finding and action plan development ?

Know Why

- Direct evidence finding instead of best guess on action plan development
- Have the vmcore provide you the opportunities to find the direct evidences ?

Know How

- Build it and test it
- Can you make a linux box with kdump enabled ?

Know What

- What is kdump ?
- How does it work ?
- Have you a clean picture about how does kdump work?

QUINN, James Brian. "The intelligent enterprise a new paradigm. The Executive," 1992, 6.4: 48-63.



Next Chapter

Courses at Grow @hp

- Course ID: 00809694 - Crash Information Capture (1.5h)
- Course ID: 00809441 - Crash(1) Commands (1.5h)
- Course ID: 00809695 - Dump Analysis Fundamentals (1.5h)
- Course ID: 00809696 - Linuxdb and web-crash (1.5h)

The Kernel books

- *Linux Kernel Development, Robert Love* ISBN: 978-0672329463
- *Understanding the Linux Kernel, Daniel P. Bovet* ISBN: 978-0596005658
- *Linux Device Drivers, Jonathan Corbet* ISBN: 978-0596005900



Thank you

